

4.yolov8 target detection

Code Path for Yolov8


```
/home/yahboom/YBAMR-COBOT-EDU-00001/soft/yolov8
```

The following figure shows the code directory file structure and folder functions.

文件名	大小	类型	修改时间	权限	用户/组
weights		文件夹	2024/06/15 17:50	drwxrwxr-x	yahboom
ultralytics		文件夹	2024/05/10 16:54	drwxrwxr-x	yahboom
torchvision		文件夹	2024/05/10 16:54	drwxrwxr-x	yahboom
runs		文件夹	2024/05/10 16:53	drwxrwxr-x	yahboom
Onnx_2_INT8		文件夹	2024/05/10 16:54	drwxrwxr-x	yahboom
models		文件夹	2024/05/10 16:54	drwxrwxr-x	yahboom
model		文件夹	2024/05/10 16:53	drwxrwxr-x	yahboom
examples		文件夹	2024/05/10 16:54	drwxrwxr-x	yahboom
docs		文件夹	2024/05/10 16:53	drwxrwxr-x	yahboom
docker		文件夹	2024/05/10 16:53	drwxrwxr-x	yahboom
dependent		文件夹	2024/06/15 17:53	drwxrwxr-x	yahboom
deep_sort		文件夹	2024/05/10 16:54	drwxrwxr-x	yahboom
coco128-images		文件夹	2024/05/10 16:54	drwxrwxr-x	yahboom
cfg		文件夹	2024/05/10 16:53	drwxrwxr-x	yahboom
.github		文件夹	2024/05/10 16:53	drwxrwxr-x	yahboom
__pycache__		文件夹	2024/06/15 11:21	drwxrwxr-x	yahboom
export.py	203 B	文件	2024/05/10 16:54	-rw-rw-r--	yahboom
CITATION.cff	612 B	文件	2024/05/10 16:53	-rw-rw-r--	yahboom
yolov8_track.py	928 B	文件	2024/05/10 16:54	-rw-rw-r--	yahboom
requirements.txt	1.2 KB	文件	2024/05/10 16:53	-rw-rw-r--	yahboom
.pre-commit-config.yaml	1.8 KB	文件	2024/05/10 16:53	-rw-rw-r--	yahboom
.gitignore	2.2 KB	文件	2024/05/10 16:54	-rw-rw-r--	yahboom
tracker_trt.py	2.9 KB	文件	2024/05/10 16:53	-rw-rw-r--	yahboom
setup.py	3.1 KB	文件	2024/05/10 16:54	-rw-rw-r--	yahboom
CONTRIBUTING.md	5.5 KB	文件	2024/05/10 16:54	-rw-rw-r--	yahboom
.DS_Store	6 KB	文件	2024/05/10 16:53	-rw-rw-r--	yahboom
multi_batch_inference.py	7.7 KB	文件	2024/06/15 18:04	-rw-rw-r--	yahboom
multi_batch_inference_track.py	9.3 KB	文件	2024/06/15 11:19	-rw-rw-r--	yahboom
mkdocs.yml	23.7 KB	文件	2024/05/10 16:53	-rw-rw-r--	yahboom
LICENSE	33.7 KB	文件	2024/05/10 16:53	-rw-rw-r--	yahboom
demo.mp4	15.5 MB	文件	2024/05/10 16:54	-rw-rw-r--	yahboom


Run the demo command of target detection and execute the following code in the terminal

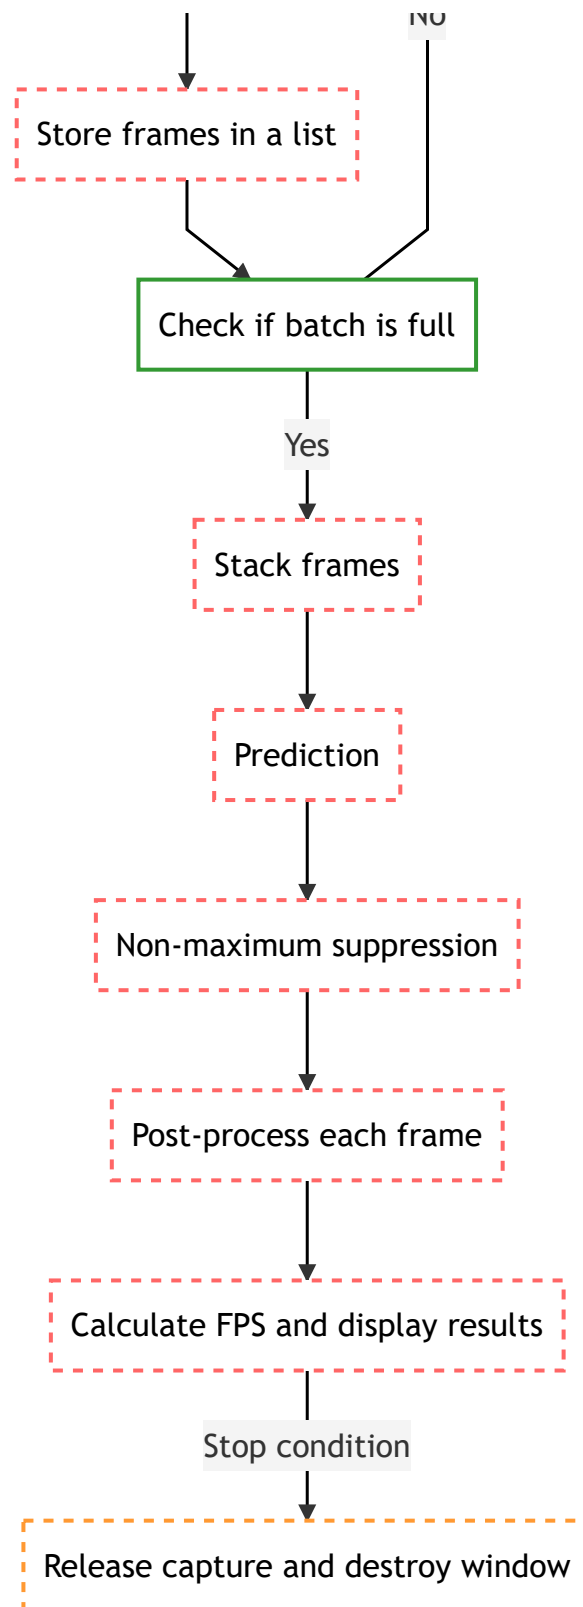
```
cd ~/YBAMR-COBOT-EDU-00001/soft/yolov8 && python multi_batch_inference.py
```



Code analysis

Code running flow chart:





Below is the code for our multi-batch object detection reasoning. The model accelerated by yolov8s' tenserrt is used.

```
import cv2
import numpy as np
from collections import OrderedDict, namedtuple
import time
import torch
```

```

from ultralytics.utils.ops import non_max_suppression, scale_boxes
import tensorrt as trt

names = ['person', 'bicycle', 'car', 'motorcycle', 'airplane', 'bus', 'train',
        'truck', 'boat', 'traffic light',
        'fire hydrant', 'stop sign', 'parking meter', 'bench', 'bird', 'cat',
        'dog', 'horse', 'sheep', 'cow',
        'elephant', 'bear', 'zebra', 'giraffe', 'backpack', 'umbrella',
        'handbag', 'tie', 'suitcase', 'frisbee',
        'skis', 'snowboard', 'sports ball', 'kite', 'baseball bat', 'baseball
glove', 'skateboard', 'surfboard',
        'tennis racket', 'bottle', 'wine glass', 'cup', 'fork', 'knife',
        'spoon', 'bowl', 'banana', 'apple',
        'sandwich', 'orange', 'broccoli', 'carrot', 'hot dog', 'pizza',
        'donut', 'cake', 'chair', 'couch',
        'potted plant', 'bed', 'dining table', 'toilet', 'tv', 'laptop',
        'mouse', 'remote', 'keyboard', 'cell phone',
        'microwave', 'oven', 'toaster', 'sink', 'refrigerator', 'book',
        'clock', 'vase', 'scissors', 'teddy bear',
        'hair drier', 'toothbrush'] # coco80 Category
colors = [[np.random.randint(0, 255) for _ in range(3)] for _ in names] ##Set
random color

def letterbox(im, new_shape=(640, 640), color=(114, 114, 114), auto=False,
scaleup=True, stride=32):
    # Resize and pad an image while satisfying stride multiple constraints
    shape = im.shape[:2] # current shape [height, width]
    if isinstance(new_shape, int):
        new_shape = (new_shape, new_shape)

    # Scale ratio (new / old)
    r = min(new_shape[0] / shape[0], new_shape[1] / shape[1])
    if not scaleup: # only scale down, do not scale up (for better val mAP)
        r = min(r, 1.0)

    # Compute padding
    new_unpad = int(round(shape[1] * r)), int(round(shape[0] * r))
    dw, dh = new_shape[1] - new_unpad[0], new_shape[0] - new_unpad[1] # wh
padding

    if auto: # minimum rectangle
        dw, dh = np.mod(dw, stride), np.mod(dh, stride) # wh padding

    dw /= 2 # divide padding into 2 sides
    dh /= 2

    if shape[::-1] != new_unpad: # resize
        im = cv2.resize(im, new_unpad, interpolation=cv2.INTER_LINEAR)
        top, bottom = int(round(dh - 0.1)), int(round(dh + 0.1))
        left, right = int(round(dw - 0.1)), int(round(dw + 0.1))
        im = cv2.copyMakeBorder(im, top, bottom, left, right, cv2.BORDER_CONSTANT,
value=color) # add border
    # print(dw,dh)
    return im, r, (dw, dh)

class TRT_engine():

```

```

def __init__(self, weight, thres=0.60, size=640, video_path='',
batch_size=3) -> None:
    self.video_path = video_path
    self.imgsz = size
    self.weight = weight
    self.iou_thres = thres
    self.batch_size = batch_size
    self.device = torch.device('cuda:0')
    self.init_engine()

def init_engine(self):
    # Infer TensorRT Engine
    self.Binding = namedtuple('Binding', ('name', 'dtype', 'shape', 'data',
'ptr'))
    self.logger = trt.Logger(trt.Logger.INFO)
    trt.init_libnvinfer_plugins(self.logger, namespace="")
    with open(self.weight, 'rb') as self.f, trt.Runtime(self.logger) as
self.runtime:
        self.model = self.runtime.deserialize_cuda_engine(self.f.read())
    self.bindings = OrderedDict()
    print(f"num binding = {self.model.num_bindings}")
    for index in range(self.model.num_bindings):
        self.name = self.model.get_binding_name(index)
        print(f"name = {self.name}")
        self.dtype = trt.nptype(self.model.get_binding_dtype(index))
        self.shape = tuple(self.model.get_binding_shape(index))
        self.data = torch.from_numpy(np.empty(self.shape,
dtype=np.dtype(self.dtype))).to(self.device)
        self.bindings[self.name] = self.Binding(self.name, self.dtype,
self.shape, self.data,
                                                    int(self.data.data_ptr()))

    self.binding_addrs = OrderedDict((n, d.ptr) for n, d in
self.bindings.items())
    self.context = self.model.create_execution_context()

def predict(self, imgs):
    self.binding_addrs['images'] = int(imgs.data_ptr())

    self.context.execute_async_v2(list(self.binding_addrs.values()),
torch.cuda.current_stream().cuda_stream)
    outputs = self.bindings['output0'].data # Compare with the results
output by onnx 393,409,425,441
    # print(outputs.shape)
    return outputs

def process(self):
    cap = cv2.VideoCapture(self.video_path)
    # cap.set(cv2.CAP_PROP_FRAME_WIDTH, 1280)
    # cap.set(cv2.CAP_PROP_FRAME_HEIGHT, 720)
    img_list = []
    square_frame_list = []
    stop = False
    while cap.isOpened() and not stop:
        ret, frame = cap.read()
        if ret:
            frame_tensor, _, dw, dh = preprocess(frame, imgsz=self.imgsz)
            img_list.append(frame_tensor)
            square_frame_list.append(frame)

```

```

        if len(img_list) == self.batch_size: # Used to store processed
frame images. When the queue is full, the images in the queue are concatenated
into a tensor.

            frames = torch.stack(img_list, 0)
            t1 = time.perf_counter()
            outputs = self.predict(frames)
            t2 = time.perf_counter()
            infer_time = (t2 - t1) / self.batch_size
            # print(outputs.shape)
            outputs = non_max_suppression(outputs, 0.25, self.iou_thres,
classes=None, agnostic=False)
            for i in range(self.batch_size):
                t3 = time.perf_counter()
                result = post_process(square_frame_list[i], outputs[i],
frames)

                t4 = time.perf_counter()
                fps = 1 / (infer_time + t4 - t3)
                cv2.putText(result, f"{fps:.2f} FPS", (15, 30),
cv2.FONT_HERSHEY_SIMPLEX, 0.7, (0, 255, 0), 2)
                cv2.imshow("result", result)
                if cv2.waitKey(1) & 0xFF == ord("q"):
                    stop = True
                    break
                img_list = []
                square_frame_list = []
            else:
                break
        cap.release()
        cv2.destroyAllWindows()

def preprocess(image, imgsz=640):
    img, ratio, (dw, dh) = letterbox(image, imgsz, stride=32,
auto=False) # When auto is FALSE, the
output image is 960*960, and when it is TRUE, it is resized proportionally.
    # The maximum length and width of the output image is 960
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB) # BGR to RGB
    img = img.transpose((2, 0, 1))
    img = np.ascontiguousarray(img)
    img = torch.from_numpy(img).to(torch.device('cuda:0')).float()
    img /= 255.0
    return img, image, dw, dh # (3,640,640)

def post_process(img, det, frames):
    """
    Draw bounding boxes on the input image.
    """
    if len(det):
        # Rescale boxes from img_size to im0 size
        det[:, :4] = scale_boxes(frames.shape[2:], det[:, :4],
img.shape).round()
        for *xyxy, conf, cls in reversed(det):
            label = f'{names[int(cls)]} {conf:.2f}'
            p1, p2 = (int(xyxy[0]), int(xyxy[1])), (int(xyxy[2]), int(xyxy[3]))
            cv2.rectangle(img, p1, p2, colors[int(cls)], thickness=4,
lineType=cv2.LINE_AA)

```

```

        cv2.putText(img, label, (p1[0], p1[1] - 5),
cv2.FONT_HERSHEY_SIMPLEX,
                    0.7, colors[int(cls)], thickness=3,
lineType=cv2.LINE_AA)
    return img

if __name__ == '__main__':
    batch_size = 4 #Set the batch size to 4
    #trt_path = f"./weights/yolov8n-cal-batch{batch_size}-int8.trt"
    trt_path = f"./weights/yolov8s.trt" ##Model Path
    trt_engine = TRT_engine(trt_path, batch_size=batch_size, thres=0.45,
size=640, video_path="demo.mp4") ##Load the video used for inference
    trt_engine.process()

```